

EC2 Instance Creation

AMI and Instance Type

▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

 Search our full catalog including 1000s of application and OS images

[AMI from catalog](#)

[Recents](#)

[Quick Start](#)

Name

Deep Learning OSS Nvidia Driver AMI GPU PyTorch 2.7 (Ubuntu 22.04)

Verified provider

Description

Release notes: <https://docs.aws.amazon.com/dlami/latest/devguide/appendix-ami-release-notes.html>

Supported EC2 instances: G4dn, G5, G6, Gr6, G6e, P4, P4de, P5, P5e, P5en, P6-B200. Release notes: <https://docs.aws.amazon.com/dlami/latest/devguide/appendix-ami-release-notes.html>

Image ID

ami-02a169a4427f8ac5b

Username

ubuntu

Catalog

Quick Start AMIs

Published

2025-07-07T10:49:21.000Z

Architecture

x86_64

Virtualization

hvm

Root device type


ebs

ENA Enabled

Yes

Boot mode

uefi-preferred



[Browse more AMIs](#)
Including AMIs from
AWS, Marketplace and
the Community

▼ Instance type [Info](#) | [Get advice](#)

Instance type

g4dn.xlarge

Family: g4dn 4 vCPU 16 GiB Memory Current generation: true On-Demand Linux base pricing: 0.526 USD per Hour
On-Demand SUSE base pricing: 0.582 USD per Hour On-Demand Ubuntu Pro base pricing: 0.533 USD per Hour
On-Demand Windows base pricing: 0.71 USD per Hour On-Demand RHEL base pricing: 0.584 USD per Hour

 All generations

[Compare instance types](#)

[Additional costs apply for AMIs with pre-installed software](#)

Access and Network Setup

▼ Key pair (login) [Info](#)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - *required*

rmisra-udacity-project-key

 [Create new key pair](#)

▼ Network settings [Info](#)

VPC - *required* [Info](#)

vpc-0b15ba35f4d67c9d5
172.31.0.0/16

(default)

Subnet [Info](#)

No preference

Auto-assign public IP [Info](#)

Enable

Additional charges apply when outside of free tier allowance.

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☒ Create security group

☐ Select existing security group

Security group name - *required*

launch-wizard-1

This security group will be added to all network interfaces. The name can't be edited after the security group is created. Max length is 255 and `[-:/()#,@!+=&:{}%*]`

Description - *required* [Info](#)

launch-wizard-1 created 2025-04-24T16:32:56.052Z

Inbound Security Group Rules

▼ Security group rule 1 (TCP, 22, 0.0.0.0/0)

[Remove](#)

Type [Info](#)

ssh

Source type [Info](#)

Anywhere

Type [Info](#)

Custom TCP

Source type [Info](#)

My IP

Type [Info](#)

Custom TCP

Source type [Info](#)

My IP

Type [Info](#)

Custom TCP

Source type [Info](#)

My IP

Protocol [Info](#)

TCP

Source [Info](#)

 Add CIDR, prefix list or security group

Protocol [Info](#)

TCP

Name [Info](#)

 Add CIDR, prefix list or security group

Protocol [Info](#)

TCP

Name [Info](#)

 Add CIDR, prefix list or security group

Protocol [Info](#)

TCP

Name [Info](#)

 Add CIDR, prefix list or security group

Port range [Info](#)

22

Description - *optional* [Info](#)

e.g. SSH for admin desktop

Port range [Info](#)

8000

Description - *optional* [Info](#)

e.g. SSH for admin desktop

Port range [Info](#)

8001

Description - *optional* [Info](#)

e.g. SSH for admin desktop

Port range [Info](#)

8002

Description - *optional* [Info](#)

e.g. SSH for admin desktop

Post Setup

Instance summary for i-0a534409c73486e94 (UdaciHeadlineEC2) [Info](#)

Updated less than a minute ago

Instance ID

[i-0a534409c73486e94](#)

IPv6 address

–

Hostname type

IP name: ip-172-31-70-50.ec2.internal

Answer private resource DNS name

IPv4 (A)

Auto-assigned IP address

[3.235.60.162](#) [Public IP]

IAM Role

–

IMDSv2

Optional

[EC2 recommends setting IMDSv2 to required](#) | [Learn more](#)

Operator

–

Public IPv4 address

[3.235.60.162](#) | [open address](#)

Instance state

Running

Private IP DNS name (IPv4 only)

[ip-172-31-70-50.ec2.internal](#)

Instance type

g4dn.xlarge

VPC ID

[vpc-05cf067a5c90c4b0b](#)

Subnet ID

[subnet-00299464f44a70cb1](#)

Instance ARN

[arn:aws:ec2:us-east-1:710271915376:instance/i-0a534409c73486e94](#)



[Connect](#)

[Instance state ▼](#)

[Actions ▼](#)

Private IPv4 addresses

[172.31.70.50](#)

Public DNS

[ec2-3-235-60-162.compute-1.amazonaws.com](#) | [open address](#)

Elastic IP addresses

–

AWS Compute Optimizer finding

[Opt-in to AWS Compute Optimizer for recommendations.](#) | [Learn more](#)

Auto Scaling Group name

–

Managed

false

Details

Status and alarms

Monitoring

Security

Networking

Storage

Tags

▼ Instance details [Info](#)

AMI ID

[ami-05ee60aff9d0a480](#)

AMI name

[Deep Learning OSS Nvidia Driver AMI GPU PyTorch 2.7 \(Ubuntu 22.04\) 2025 0602](#)

Stop protection

Disabled

Instance reboot migration

Default (On)

Stop-hibernate behavior

Disabled

State transition reason

Monitoring

disabled

Allowed image

–

Launch time

[Tue Jul 15 2025 07:37:15 GMT-0700 \(Pacific Daylight Time\)](#) (1 minute)

Instance auto-recovery

Default

AMI Launch index

0

Credit specification

Platform details

[Linux/UNIX](#)

Termination protection

Disabled

AMI location

[amazon/Deep Learning OSS Nvidia Driver AMI GPU PyTorch 2.7 \(Ubuntu 22.04\) 20250602](#)

Lifecycle

normal

Key pair assigned at launch

[vockey](#)

Kernel ID

Welcome to Ubuntu 22.04.5 LTS (GNU/Linux 6.8.0-1030-aws x86_64)

- * Documentation: <https://help.ubuntu.com>
- * Management: <https://landscape.canonical.com>
- * Support: <https://ubuntu.com/pro>

System information as of Sat Jul 12 19:42:45 UTC 2025

System load:	0.28	Processes:	164
Usage of /:	63.2% of 145.19GB	Users logged in:	1
Memory usage:	4%	IPv4 address for ens5:	172.31.68.64
Swap usage:	0%		

=> There are 2 zombie processes.

- * Ubuntu Pro delivers the most comprehensive open source security and compliance features.

<https://ubuntu.com/aws/pro>

Expanded Security Maintenance for Applications is not enabled.

15 updates can be applied immediately.

To see these additional updates run: `apt list --upgradable`

3 additional security updates can be applied with ESM Apps.

Learn more about enabling ESM Apps service at <https://ubuntu.com/esm>

The list of available updates is more than a week old.

To check for new updates run: `sudo apt update`

New release '24.04.2 LTS' available.

Run 'do-release-upgrade' to upgrade to it.

3 updates could not be installed automatically. For more details,
see `/var/log/unattended-upgrades/unattended-upgrades.log`

=====

AMI Name: Deep Learning OSS Nvidia Driver AMI GPU PyTorch 2.7 (Ubuntu 22.04)
Supported EC2 instances: G4dn, G5, G6, Gr6, G6e, P4, P4de, P5, P5e, P5en, P6-B200
* To activate pre-built pytorch environment, run: 'source /opt/pytorch/bin/activate '
NVIDIA driver version: 570.133.20
CUDA versions available: cuda-12.8
Default CUDA version is 12.8

Release notes: <https://docs.aws.amazon.com/dlami/latest/devguide/appendix-ami-release-notes.html>

AWS Deep Learning AMI Homepage: <https://aws.amazon.com/machine-learning/amis/>

Developer Guide and Release Notes: <https://docs.aws.amazon.com/dlami/latest/devguide/what-is-dlami.html>

Support: <https://forums.aws.amazon.com/forum.jspa?forumID=263>

For a fully managed experience, check out Amazon SageMaker at <https://aws.amazon.com/sagemaker>

=====

Last login: Sat Jul 12 19:32:36 2025 from 98.51.210.84

ubuntu@ip-172-31-68-64:~\$

Access EC2 Instance

```
ubuntu@ip-172-31-72-27:~$ curl -s https://pyenv.run | bash
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload    Total   Spent    Left   Speed
100    270    100    270     0     0   1953      0 --:--:-- --:--:-- --:--:--   1956
```

WARNING: Can not proceed with installation. Kindly remove the '/home/ubuntu/.pyenv' directory first.

```
ubuntu@ip-172-31-72-27:~$ echo 'export PYENV_ROOT="$HOME/.pyenv"' >> ~/.bashrc
ubuntu@ip-172-31-72-27:~$ echo 'export PATH="$PYENV_ROOT/bin:$PATH"' >> ~/.bashrc
ubuntu@ip-172-31-72-27:~$ echo 'eval "$(pyenv init --path)"' >> ~/.bashrc
ubuntu@ip-172-31-72-27:~$ echo 'eval "$(pyenv init -)"' >> ~/.bashrc
ubuntu@ip-172-31-72-27:~$ exec "$SHELL"
```

```
ubuntu@ip-172-31-72-27:~$ pyenv install 3.10.14
pyenv: /home/ubuntu/.pyenv/versions/3.10.14 already exists
continue with installation? (y/N) ubuntu@ip-172-31-72-27:~$
```

```
ubuntu@ip-172-31-72-27:~$
```

```
ubuntu@ip-172-31-72-27:~$
```

```
ubuntu@ip-172-31-72-27:~$ python --version
Python 3.10.14
```

```
ubuntu@ip-172-31-72-27:~$ Python 3.10.14
Command 'Python' not found, did you mean:
  command 'jython' from deb jython (2.7.2+repack1-4)
Try: sudo apt install <deb name>
```

```
ubuntu@ip-172-31-72-27:~$
```

```
ubuntu@ip-172-31-72-27:~$ python -m venv ~/trt-llm-env
ubuntu@ip-172-31-72-27:~$ source ~/trt-llm-env/bin/activate
(trt-llm-env) ubuntu@ip-172-31-72-27:~$
(trt-llm-env) ubuntu@ip-172-31-72-27:~$ sudo apt-get update
sudo apt-get install git-lfs -y
```

Python & Virtual Environment Setup

```
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy InRelease
Get:2 http://security.ubuntu.com/ubuntu jammy-security InRelease [129 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates InRelease [128 kB]
Get:4 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-backports InRelease [127 kB]
Hit:5 https://nvidia.github.io/libnvidia-container/stable/deb/amd64 InRelease
Hit:6 https://apt.corretto.aws stable InRelease
Get:7 https://download.docker.com/linux/ubuntu jammy InRelease [48.8 kB]
Get:8 https://fsx-lustre-client-repo.s3.amazonaws.com/ubuntu jammy InRelease [3823 B]
Get:9 https://developer.download.nvidia.com/compute/cuda/repos/ubuntu2204/x86_64 InRelease [1581 B]
Get:10 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/main amd64 Packages [2723 kB]
Get:11 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/main Translation-en [434 kB]
Get:12 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/restricted amd64 Packages [3898 kB]
Get:13 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/restricted Translation-en [698 kB]
Get:14 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/universe amd64 Packages [1221 kB]
Get:15 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/universe Translation-en [301 kB]
Get:16 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/multiverse amd64 Packages [47.1 kB]
Get:17 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-updates/multiverse Translation-en [12.0 kB]
Get:18 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-backports/main amd64 Packages [68.8 kB]
Get:19 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-backports/universe amd64 Packages [30.0 kB]
Get:20 http://us-east-1.ec2.archive.ubuntu.com/ubuntu jammy-backports/universe Translation-en [16.5 kB]
```



```
(trt-llm-env) ubuntu@ip-172-31-72-27:~$ git clone https://huggingface.co/gpt2 /home/ubuntu/models_src/gpt2
```

```
Cloning into '/home/ubuntu/models_src/gpt2'...
```

```
remote: Enumerating objects: 87, done.
```

```
remote: Counting objects: 100% (3/3), done.
```

```
remote: Compressing objects: 100% (2/2), done.
```

```
remote: Total 87 (delta 0), reused 0 (delta 0), pack-reused 84 (from 1)
```

```
Unpacking objects: 100% (87/87), 1.65 MiB | 8.50 MiB/s, done.
```

```
Filtering content: 100% (11/11), 5.23 GiB | 77.46 MiB/s, done.
```

```
(trt-llm-env) ubuntu@ip-172-31-72-27:~$ cd /home/ubuntu
```

```
(trt-llm-env) ubuntu@ip-172-31-72-27:~$ git clone https://github.com/NVIDIA/TensorRT-LLM.git
```

```
Cloning into 'TensorRT-LLM'...
```

```
remote: Enumerating objects: 95583, done.
```

```
remote: Counting objects: 100% (227/227), done.
```

```
remote: Compressing objects: 100% (146/146), done.
```

```
remote: Total 95583 (delta 145), reused 82 (delta 81), pack-reused 95356 (from 2)
```

```
Receiving objects: 100% (95583/95583), 1.59 GiB | 46.42 MiB/s, done.
```

```
Resolving deltas: 100% (69921/69921), done.
```

```
Updating files: 100% (5877/5877), done.
```

```
Filtering content: 100% (2377/2377), 1.79 GiB | 60.28 MiB/s, done.
```

```
(trt-llm-env) ubuntu@ip-172-31-72-27:~$ sudo apt-get -y install libopenmpi-dev && pip3 install --upgrade pip setuptools && pip3 install tensorrt_llm
```

```
Reading package lists... Done
```

```
Building dependency tree... Done
```

```
Reading state information... Done
```

```
libopenmpi-dev is already the newest version (4.1.2-2ubuntu1).
```

```
0 upgraded, 0 newly installed, 0 to remove and 49 not upgraded.
```

```
Requirement already satisfied: pip in ./trt-llm-env/lib/python3.10/site-packages (25.1.1)
```

```
Requirement already satisfied: setuptools in ./trt-llm-env/lib/python3.10/site-packages (79.0.1)
```

```
Collecting setuptools
```

```
Using cached setuptools-80.9.0-py3-none-any.whl.metadata (6.6 kB)
```

```
Using cached setuptools-80.9.0-py3-none-any.whl (1.2 MB)
```

```
Installing collected packages: setuptools
```

```
Attempting uninstall: setuptools
```

```
Found existing installation: setuptools 79.0.1
```

```
Uninstalling setuptools-79.0.1:
```

```
Successfully uninstalled setuptools-79.0.1
```

```
Successfully installed setuptools-80.9.0
```

```
Requirement already satisfied: tensorrt_llm in ./trt-llm-env/lib/python3.10/site-packages (0.20.0)
```

```
Requirement already satisfied: accelerate==0.25.0 in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (1.8.1)
```

```
Requirement already satisfied: build in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (1.2.2.post1)
```

```
Requirement already satisfied: colored in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (2.3.0)
```

```
Requirement already satisfied: cuda-python in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (12.9.0)
```

```
Requirement already satisfied: diffusers>=0.27.0 in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (0.34.0)
```

```
Requirement already satisfied: lark in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (1.2.2)
```

```
Requirement already satisfied: mpi4py in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (4.1.0)
```

```
Requirement already satisfied: numpy<2 in ./trt-llm-env/lib/python3.10/site-packages (from tensorrt_llm) (1.26.4)
```

GPT2 and TensorRT-LLM Setup

Model Checkpoint Conversion and Engine Build

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ python3 /home/ubuntu/TensorRT-LLM/examples/models/core/gpt/convert_checkpoint.py --model_dir /home/ubuntu/models_src/gpt2 --output_dir /home/ubuntu/models_trt/gpt2/tllm_checkpoint --dtype float16'
```

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$
```

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$
```

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ tree /home/ubuntu/models_trt/gpt2/tllm_checkpoint
```

```
/home/ubuntu/models_trt/gpt2/tllm_checkpoint
```

```
├── config.json
└── rank0.safetensors
```

```
0 directories, 2 files
```

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ trtllm-build --checkpoint_dir /home/ubuntu/models_trt/gpt2/tllm_checkpoint --output_dir /home/ubuntu/models_trt/gpt2/trt_engine_fp16_t4 --gemm_plugin float16 --remove_input_padding enable --logits_dtype float16 --max_batch_size 1 --max_seq_len 1024 --workers 1 --log_level info --context_fmha disable
```

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ tree /home/ubuntu/models_trt/gpt2/trt_engine_fp16_t4
```

```
/home/ubuntu/models_trt/gpt2/trt_engine_fp16_t4
```

```
├── config.json
└── rank0.engine
```

```
0 directories, 2 files
```

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ cp TensorRT-LLM/triton_backend/all_models/gpt/tensorrt_llm/config.pbtxt /home/ubuntu/triton_model_repo/gpt2/config.pbtxt
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ MODEL_FOLDER=/home/ubuntu/triton_model_repo
FILL_TEMPLATE_SCRIPT=/home/ubuntu/TensorRT-LLM/triton_backend/tools/fill_template.py
ENGINE_DIR=/home/ubuntu/triton_model_repo/gpt2/1
(trt-llm-env) ubuntu@ip-172-31-68-64:~$
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ python3 ${FILL_TEMPLATE_SCRIPT} ${MODEL_FOLDER}/gpt2/config.pbtxt engine_dir:${ENGINE_DIR},name:gpt2
```

```
name: "tensorrt_llm"
backend: "python"
max_batch_size: 1024
```

```
# # Uncomment this for dynamic_batching
# dynamic_batching {
#     max_queue_delay_microseconds: 50000
# }
```

```
input {
  {
    name: "input_ids"
    data_type: TYPE_INT32
    dims: [ -1 ]
  },
  {
    name: "input_lengths"
    data_type: TYPE_INT32
    dims: [ 1 ]
    reshape: { shape: [ ] }
  },
  {
    name: "request_output_len"
    data_type: TYPE_INT32
    dims: [ -1 ]
  },
  {
    name: "end_id"
    data_type: TYPE_INT32
    dims: [ 1 ]
    reshape: { shape: [ ] }
  },
  {
    name: "pad_id"
    data_type: TYPE_INT32
    dims: [ 1 ]
    reshape: { shape: [ ] }
  },
  {
    name: "beam_width"
    data_type: TYPE_INT32
    dims: [ 1 ]
    reshape: { shape: [ ] }
    optional: true
  },
  {
    name: "temperature"
    data_type: TYPE_FP32
```

Config.pbtxt Creation


```

cp /home/ubuntu/models_trt/gpt2/trt_engine_fp16_t4 /rank0.engine
/home/ubuntu/triton_model_repo/gpt2/ 1/rank0.engine

cp /home/ubuntu/models_trt/gpt2/trt_engine_fp16_t4 /config.json /home/ubuntu/triton_model_repo/gpt2/ 1/config.json

cp TensorRT-LLM/triton_backend/all_models/gpt/tensorrt_llm/1/model.py
/home/ubuntu/triton_model_repo/gpt2/1/model.py

cp /home/ubuntu/models_src/gpt2/tokenizer_config.json
/home/ubuntu/triton_model_repo/gpt2/1/tokenizer/tokenizer_config.json

cp /home/ubuntu/models_src/gpt2/tokenizer.json /home/ubuntu/triton_model_repo/gpt2/1/tokenizer/tokenizer.json

cp /home/ubuntu/models_src/gpt2/vocab.json /home/ubuntu/triton_model_repo/gpt2/1/tokenizer/vocab.json

cp /home/ubuntu/models_src/gpt2/merges.txt /home/ubuntu/triton_model_repo/gpt2/1/tokenizer/merges.txt

```

Copy Model Files & Check Repo Structure

```

(trt-llm-env) ubuntu@ip-172-31-68-64:~$ tree /home/ubuntu/triton_model_repo
/home/ubuntu/triton_model_repo
├── gpt2
│   ├── 1
│   │   ├── pycache_
│   │   │   └── model.cpython-312.pyc
│   │   ├── config.json
│   │   ├── model.py
│   │   ├── rank0.engine
│   │   └── tokenizer
│   │       ├── merges.txt
│   │       ├── tokenizer.json
│   │       ├── tokenizer_config.json
│   │       └── vocab.json
│   └── config.pbtxt
4 directories, 9 files
(trt-llm-env) ubuntu@ip-172-31-68-64:~$

```

Start Triton Server

```
(trt-llm-env) ubuntu@ip-172-31-68-64:~$ docker run --rm -it --net host --gpus all \
--shm-size=2g --ulimit memlock=-1 --ulimit stack=67108864 \
-v /home/ubuntu/triton_model_repo:/home/ubuntu/triton_model_repo \
nvcr.io/nvidia/tritonserver:25.06-trtllm-python-py3 \
tritonserver --model-repository=/home/ubuntu/triton_model_repo
```

```
=====
== Triton Inference Server ==
=====
```

NVIDIA Release 25.06 (build 179868725)
Triton Server Version 2.59.0

Copyright (c) 2018-2025, NVIDIA CORPORATION & AFFILIATES. All rights reserved.

Various files include modifications (c) NVIDIA CORPORATION & AFFILIATES. All rights reserved.

GOVERNING TERMS: The software and materials are governed by the NVIDIA Software License Agreement
(found at <https://www.nvidia.com/en-us/agreements/enterprise-software/nvidia-software-license-agreement/>)
and the Product-Specific Terms for NVIDIA AI Products
(found at <https://www.nvidia.com/en-us/agreements/enterprise-software/product-specific-terms-for-ai-products/>).

NOTE: CUDA Forward Compatibility mode ENABLED.

Using CUDA 12.9 driver version 575.51.02 with kernel driver version 535.247.01.
See <https://docs.nvidia.com/deploy/cuda-compatibility/> for details.

Triton Server Status

```
[TensorRT-LLM] TensorRT-LLM version: 0.20.0
I0714 22:24:39.321789 1 model_lifecycle.cc:849] "successfully loaded 'gpt2'"
I0714 22:24:39.321920 1 server.cc:611]
+-----+
| Repository Agent | Path |
+-----+
+-----+

I0714 22:24:39.321974 1 server.cc:638]
+-----+
| Backend | Path | Config |
+-----+
+-----+

| python | /opt/tritonserver/backends/python/libtriton_python.so | {"cmdline":{"auto-complete-config":"true","backend-directory":"/opt/tritonserver/backends","min-compute-capability":"6.000000","default-max-batch-size":"4"}} |
+-----+
+-----+

I0714 22:24:39.322074 1 server.cc:681]
+-----+
| Model | Version | Status |
+-----+
| gpt2 | 1 | READY |
+-----+
+-----+

I0714 22:24:39.401656 1 metrics.cc:890] "Collecting metrics for GPU 0: Tesla T4"
I0714 22:24:39.406242 1 metrics.cc:783] "Collecting CPU metrics"
I0714 22:24:39.406431 1 tritonserver.cc:2598]
+-----+
+-----+

| Option | Value |
+-----+
+-----+

| server_id | triton |
+-----+
| server_version | 2.59.0 |
+-----+
| server_extensions | classification sequence model_repository model_repository(unload_dependents) schedule_policy model_configuration system_shared_memory cuda_shared_memory binary_tensor_data parameters statistics trace logging |
+-----+
| model_repository_path[0] | /home/ubuntu/triton_model_repo |
+-----+
| model_control_mode | MODE_NONE |
+-----+
| strict_model_config | 0 |
+-----+
| model_config_name | |
+-----+
| rate_limit | OFF |
+-----+
| pinned_memory_pool_byte_size | 268435456 |
+-----+
| cuda_memory_pool_byte_size{0} | 67108864 |
+-----+
| min_supported_compute_capability | 6.0 |
+-----+
| strict_readiness | 1 |
+-----+
| exit_timeout | 30 |
+-----+
| cache_enabled | 0 |
+-----+
+-----+

I0714 22:24:39.445471 1 grpc_server.cc:2562] "Started GRPCInferenceService at 0.0.0.0:8001"
I0714 22:24:39.446462 1 http_server.cc:4832] "Started HTTPService at 0.0.0.0:8000"
I0714 22:24:39.489409 1 http_server.cc:358] "Started Metrics Service at 0.0.0.0:8002"
```